

# Week 11


## Learnings

### Prometheus

I set up Prometheus using this tutorial: <https://devopscube.com/setup-prometheus-monitoring-on-kubernetes/>. Originally I was letting a package manager configure everything for prometheus for me, but when I ran into issues, it was hard to diagnose since I had no clue what I was touching.

After setting up Prometheus I need to set up kube-state-metrics so I can get information (like how many pods are running in a namespace). I followed this tutorial: <https://devopscube.com/setup-kube-state-metrics/>. This allowed me to view the amount of pods running per namespace using the following PromQL command in the Prometheus UI.

```
count(kube_pod_info{namespace="app"}) by (created_by_name)
```



The following command will show the CPU utilization for the backend pods in the app namespace

```
sum by (namespace, pod) (
  rate(container_cpu_usage_seconds_total{namespace="app",pod=~"backend-.*"}[5m])
)
```

# Kubernetes

## Namespaces

Provide a mechanism for isolating groups of resources within a single cluster.

Kubernetes starts with four initial namespaces:

1. *default*: So you can start using the cluster without first creating a namespace.
2. *kube-node-lease*: Holds the “Lease” objects which allow the kubelet to send heartbeats so that the control plan can detect node failure.
3. *kube-public*: ?
4. *kube-system*: For objects created by the Kubernetes system. For example, this is where the metrics-server goes when you initialize it so that prometheus can use it.

To create a namespace, you can run the following command to define the kubernetes namespaces in the yaml files.

```
kubectl apply -f kubernetes/namespaces
```

Then, to use the namespace, first you must create a context. The following command creates a context called app and sets it to map to the “app” namespace from the folder

kubernetes/namespaces .

```
kubectl config set-context app --namespace=app \
  --cluster=minikube \
  --user=minikube
```

Finally, to set the current context to the app name space (so you can deploy your pods from there), you can run the following command:

```
kubectl config use-context app
```

Now you can do all the normal things you would with a namespace like applying deployments, services and autoscalers.

## Config Maps


Kubernetes config maps allow you to externalize some configuration so that you don't need to redeploy each time you need to change some configuration.

## Jobs

Creates one or more Pods and will continue to retry execution of the Pods until a specified number of them successfully terminate.

# Problems

**Grafana is taking an extremely long time to load.**



```
E1121 06:32:18.150974 11667 portforward.go:370] error creating forwarding stream for port 3000 -> 3000: Timeout occurred
E1121 06:32:24.247463 11667 portforward.go:347] error creating error stream for port 3000 -> 3000: Timeout occurred
E1121 06:32:24.247550 11667 portforward.go:347] error creating error stream for port 3000 -> 3000: Timeout occurred
Handling connection for 3000
Handling connection for 3000
E1121 06:32:24.259170 11667 portforward.go:347] error creating error stream for port 3000 -> 3000: Timeout occurred
E1121 06:32:24.259197 11667 portforward.go:347] error creating error stream for port 3000 -> 3000: Timeout occurred
Handling connection for 3000
E1121 06:32:24.259993 11667 portforward.go:347] error creating error stream for port 3000 -> 3000: Timeout occurred
Handling connection for 3000
Handling connection for 3000
```

It looks like Prometheus is getting restarted constantly because it's being deemed unhealthy? This is not a problem I've had when starting minikube previously

prometheus-grafana-6fdbff5c9-djvr:17999fc9a1814b27	Unhealthy	Readiness probe failed: Get "http://10.244.0.135:3000/api/health": dial tcp 10.244.0.135:3000: connect: connection refused	kubelet minikube	spec.containers(grafana)	2	7 minutes ago	7 minutes ago
--	-----------	--	------------------	--------------------------	---	---------------	---------------

## Solution:

I restarted the port-forward, waiting until the grafana pods were actually up.

# Additional

## Commands for memory

### Restarting a deployment

```
kubectl rollout restart deployment <deployment_name>
```

### Unapplying an applied yaml file

```
kubectl delete -f <file_location>
```